

**COURSE:** DSA4050

**COURSE TITLE:** Deep Learning for Computer Vision

**INSTRUCTOR:** DR. EDWARD OMBUI

## **Lesson Notes: Semantic Segmentation and Instance Segmentation**

---

### **Week 6 Overview**

This week, we explore **segmentation tasks in computer vision**, where the goal is to classify every pixel in an image. Unlike traditional image classification, which assigns a single label to an image, segmentation tasks aim to **identify objects at the pixel level**.

We will focus on two key segmentation techniques:

1. **Semantic Segmentation:** Classifies all pixels belonging to the same category as one.
  - Example: Identifying **all cars** in an image but not distinguishing between individual cars.
2. **Instance Segmentation:** Not only classifies objects but also **distinguishes between individual instances** of the same class.
  - Example: Separating **one car from another** in the same image.

You will also study **popular segmentation architectures** such as:

- **Fully Convolutional Networks (FCNs):** Transform traditional CNNs into pixel-level classifiers.
- **U-Net:** An encoder-decoder network widely used for **biomedical image segmentation**.

By the end of this week, you will be able to **differentiate between segmentation types, understand segmentation models, and implement segmentation tasks using FCNs and U-Net**.

## **1. Understanding Semantic and Instance Segmentation**

### **What is Semantic Segmentation?**

Semantic segmentation assigns each pixel to a **category** (e.g., car, road, pedestrian). All objects of the same category are given the same label, without distinguishing between different instances.

### Example Use Cases:

- ✓ **Autonomous Driving:** Identifying road, lanes, vehicles, and pedestrians.
- ✓ **Medical Imaging:** Segmenting organs, tumors, or abnormalities in MRI scans.
- ✓ **Aerial Image Analysis:** Identifying land, water, and buildings in satellite images.

### Example Output:

- Input: Image with multiple cars.
- Output: A mask where all car pixels are marked in **one color**.

## What is Instance Segmentation?

Instance segmentation goes beyond semantic segmentation by **differentiating each object instance** within the same class.

### Example Use Cases:

- ✓ **Self-Driving Cars:** Detecting and tracking individual vehicles and pedestrians.
- ✓ **Retail and Inventory Management:** Counting products and monitoring stock.
- ✓ **Healthcare:** Identifying different cell types in microscopy images.

### Example Output:

- Input: Image with multiple people.
- Output: A mask where each person has a **unique color** to distinguish between individuals.

☐ **Reference:** He et al. (2017). *Mask R-CNN: Instance Segmentation Framework*.

## 2. Segmentation Architectures: FCNs and U-Net

### 2.1 Fully Convolutional Networks (FCNs)

★ **Key Idea:** Transform a traditional CNN into a fully convolutional model where the output is a **pixel-wise** prediction.

#### ✓ **Strengths:**

- Efficient for dense predictions.
- Can adapt classification models for segmentation.

### ✗ Limitations:

- Struggles with preserving fine-grained details.
- May produce blurry segmentation boundaries.

### 🔧 Implementation of FCN using TensorFlow

```
import tensorflow as tf
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Conv2D, Conv2DTranspose, Input

# Define FCN model
input_layer = Input(shape=(256, 256, 3))
conv1 = Conv2D(64, (3, 3), activation='relu', padding='same')(input_layer)
conv2 = Conv2D(128, (3, 3), activation='relu', padding='same')(conv1)
conv3 = Conv2DTranspose(64, (3, 3), strides=(2, 2), activation='relu',
padding='same')(conv2)
output_layer = Conv2D(1, (1, 1), activation='sigmoid')(conv3)

model = Model(inputs=input_layer, outputs=output_layer)
model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])
model.summary()
```

📖 **Reference:** Long et al. (2015). *Fully Convolutional Networks for Semantic Segmentation*.

## 2.2 U-Net: Encoder-Decoder Architecture

★ **Key Idea:** Uses an **encoder-decoder** structure with skip connections to capture both **global context** and **fine details**.

### ✓ Strengths:

- Works well for **biomedical image segmentation**.
- Maintains high accuracy with limited data.

### ✗ Limitations:

- Computationally intensive for large images.
- Requires data augmentation for better generalization.

## 🔧 Implementation of U-Net using TensorFlow

```
from tensorflow.keras.layers import MaxPooling2D, concatenate

def unet_model(input_shape=(256, 256, 3)):
    inputs = Input(shape=input_shape)

    # Encoder
    conv1 = Conv2D(64, (3, 3), activation='relu', padding='same')(inputs)
    pool1 = MaxPooling2D((2, 2))(conv1)

    conv2 = Conv2D(128, (3, 3), activation='relu', padding='same')(pool1)
    pool2 = MaxPooling2D((2, 2))(conv2)

    # Decoder
    up1 = Conv2DTranspose(64, (3, 3), strides=(2, 2), activation='relu',
padding='same')(pool2)
    merge1 = concatenate([conv1, up1], axis=3)
    output_layer = Conv2D(1, (1, 1), activation='sigmoid')(merge1)

    model = Model(inputs=inputs, outputs=output_layer)
    model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])
    return model

model = unet_model()
model.summary()
```

❏ **Reference:** Ronneberger et al. (2015). *U-Net: Convolutional Networks for Biomedical Image Segmentation*.

## 3. Training and Evaluating a Segmentation Model

### 3.1 Dataset Preparation

1. **Choose a Dataset:** Common datasets include **COCO**, **PASCAL VOC**, **Cityscapes**, or a **custom medical dataset**.
2. **Preprocess Images:** Resize images, normalize pixel values, and apply data augmentation.

### 3.2 Training the Model

```
history = model.fit(train_images, train_masks, epochs=10,
validation_data=(val_images, val_masks))
```

### 3.3 Evaluation Metrics

Metric	Description
<b>IoU (Intersection over Union)</b>	Measures overlap between predicted and ground truth masks.
<b>Dice Coefficient</b>	Measures similarity between two sets (higher is better).
<b>Pixel Accuracy</b>	Percentage of correctly classified pixels.

## 4. Learning Outcomes

By the end of this session, you should be able to:

- ✓ **Differentiate** between **semantic segmentation** and **instance segmentation**.
- ✓ **Describe and implement Fully Convolutional Networks (FCNs)** and **U-Net**.
- ✓ **Train and test a segmentation model** using a dataset.

## 5. Laboratory Assignment

### Task:

**Implement Fully Convolutional Networks (FCNs) and U-Net for segmentation tasks.**

### Steps:

1. Select a dataset (e.g., medical, street scene segmentation).
2. Implement both **FCN** and **U-Net architectures**.
3. Train the models using **IoU** and **Dice Coefficient** as metrics.
4. Compare model performance.

### ★ Deliverables:

- Trained segmentation models.
- Evaluation report comparing **FCN vs. U-Net** performance.
- Sample segmentation results.

📖 **Reference:** Ronneberger et al. (2015). *U-Net for Biomedical Image Segmentation*.

## 6. Conclusion

- ✓ **Semantic segmentation** classifies **all pixels** in an image.
- ✓ **Instance segmentation** differentiates **individual objects** within the same category.
- ✓ **FCNs and U-Net** are widely used **segmentation architectures**.
- ✓ **Segmentation models are evaluated using IoU, Dice Coefficient, and Pixel Accuracy.**

### ✦ Next Steps:

- Experiment with **custom datasets**.
- Use **transfer learning** to improve segmentation accuracy.
- Apply segmentation to **real-world applications** like medical imaging and self-driving cars.